


GPT-4o vs Claude 3.5 Sonnet — Cost Comparison

Prices per 1 million tokens. Source: OpenAI & Anthropic official pricing pages, May 2026.



■ GPT-4o ■ Claude 3.5 Sonnet

Red % = how much more Claude costs vs GPT-4o for same volume

The Numbers at a Glance

	GPT-4o	Claude 3.5 Sonnet	Difference
Input (per 1M tokens)	\$2.50	\$3.00	Claude costs \$0.50 more (+20%)
Output (per 1M tokens)	\$10.00	\$15.00	Claude costs \$5.00 more (+50%)
Batch API Input	\$1.25	\$1.50	50% discount off standard for both
Batch API Output	\$5.00	\$7.50	50% discount off standard for both
Context Window	128k tokens	200k tokens	Claude fits 56% more in one request

Batch API = send a large job, get results within 24 hours. Both providers give a flat 50% discount for this.

What Does This Cost in Practice?

Three realistic scenarios — calculated using confirmed prices.

Three real-world scenarios — what you actually pay

Small app (100k requests/mo) Each request: 500 input + 200 output tokens GPT-4o \$325.00 Claude costs 38% more Claude 3.5 Sonnet \$450.00	Medium app (1M requests/mo) Each request: 800 input + 300 output tokens GPT-4o \$5,000.00 Claude costs 38% more Claude 3.5 Sonnet \$6,900.00	Document processing (1-time) 1M input tokens (a ~750-page book) + 50k output tokens GPT-4o \$3.00 Claude costs 25% more Claude 3.5 Sonnet \$3.75
---	--	--

Maths Behind the Scenarios

Scenario 1 — Small app (100,000 requests/month)

Each request uses 500 input tokens + 200 output tokens. That gives us $100,000 \times 500 = 50$ million input tokens, and $100,000 \times 200 = 20$ million output tokens per month.

	Tokens used	GPT-4o cost	Claude 3.5 cost
Input	50M tokens	$50 \times \$2.50 = \125.00	$50 \times \$3.00 = \150.00
Output	20M tokens	$20 \times \$10.00 = \200.00	$20 \times \$15.00 = \300.00
Total / month	70M tokens	\$325.00	\$450.00

Scenario 2 — Medium app (1,000,000 requests/month)

Each request uses 800 input tokens + 300 output tokens. That gives us 800 million input tokens and 300 million output tokens per month.

	Tokens used	GPT-4o cost	Claude 3.5 cost
Input	800M tokens	$800 \times \$2.50 = \$2,000.00$	$800 \times \$3.00 = \$2,400.00$
Output	300M tokens	$300 \times \$10.00 = \$3,000.00$	$300 \times \$15.00 = \$4,500.00$
Total / month	1,100M tokens	\$5,000.00	\$6,900.00

Scenario 3 — One-time document processing

Sending 1 million input tokens (roughly a 750-page book) and receiving 50,000 output tokens (a detailed summary).

	Tokens used	GPT-4o cost	Claude 3.5 cost
Input	1M tokens	$1 \times \$2.50 = \2.50	$1 \times \$3.00 = \3.00
Output	50k tokens	$0.05 \times \$10.00 = \0.50	$0.05 \times \$15.00 = \0.75
Total (one-time)	1.05M tokens	\$3.00	\$3.75

When to Pick Which

Pick GPT-4o if...	Pick Claude 3.5 Sonnet if...
Cost is your top priority and you are running high-volume output-heavy workloads (output is where the price gap is biggest: \$10 vs \$15)	You need a longer context window — Claude fits 200k tokens in one call vs GPT-4o's 128k. That is ~150 pages vs ~96 pages of text in a single request.
You are already deep in the OpenAI ecosystem (Assistants API, function calling, etc.)	Your task involves long documents, large codebases, or detailed reasoning where the extra context window removes the need to split and chunk your data.
You want to use the Batch API for async jobs — GPT-4o batch input (\$1.25) is slightly cheaper than Claude batch input (\$1.50)	The price premium is acceptable for your use case — \$3.75 vs \$3.00 for a 750-page document is a \$0.75 difference per job.

Bottom line: GPT-4o is cheaper — especially on output tokens (\$10 vs \$15). Claude 3.5 Sonnet costs 20% more on input and 50% more on output, but gives you a 200k context window vs 128k. If price is your only factor, GPT-4o wins. If you need to fit more into a single call, Claude's larger window may save you the cost of chunking.

Prices verified May 2026. Always check openai.com/api/pricing and anthropic.com/pricing before making production cost decisions — prices change.