

LLM Model Providers

Flagship Models & Pricing · Month 1, Week 1, Task 2 · Prices as of May 2026

There are five major players in the LLM landscape you need to know as an AI engineer. Three are **closed-source** providers you call via a paid API (OpenAI, Anthropic, Google). One releases **open-source weights** you can run yourself or via a third-party host (Meta). One sits in between — open weights for smaller models, paid API for its flagship (Mistral). Understanding who does what, and what it costs, is a core engineering decision you will make on every project.

The Five Providers at a Glance

OpenAI Closed-source API	Flagship: GPT-4o · GPT-4.1 · o3 Pricing (input): \$0.40 – \$10 / 1M tokens CLOSED API
Anthropic Closed-source API	Flagship: Claude Opus 4.7 · Sonnet 4.6 · Haiku 4.5 Pricing (input): \$1 – \$25 / 1M tokens CLOSED API
Google Closed-source API	Flagship: Gemini 2.5 Pro · Gemini 2.5 Flash Pricing (input): \$0.05 – \$1.25 / 1M tokens CLOSED API
Meta Open-source weights	Flagship: Llama 3.3 70B · Llama 3.1 8B Pricing (input): \$0.02 – \$0.40 / 1M tokens (via host) OPEN WEIGHTS
Mistral Open + closed API	Flagship: Mistral Large 3 · Medium 3 · Devstral 2 Pricing (input): \$0.40 – \$2 / 1M tokens OPEN WEIGHTS

Input Token Pricing Spectrum (per 1M tokens)



Figure — Input token cost comparison across flagship and mid-tier models. OSS = open-source, hosted via third-party providers like Together AI, Groq, or AWS Bedrock. Output tokens typically cost 3–5× more than input tokens.

Who they are

OpenAI created the modern LLM industry with GPT-3 and later ChatGPT. They remain the largest AI company by API adoption and have the broadest ecosystem — most third-party tools, tutorials, and integrations default to OpenAI first. They offer two model families: the **GPT series** (general-purpose, optimised for instruction following and coding) and the **o-series** (reasoning models that 'think before answering', better for complex maths and logic).

Model lineup & pricing

Model	Best for	Input / 1M tokens	Output / 1M tokens	Context window
GPT-4o	General tasks, vision, low latency	\$2.50	\$10.00	128K tokens
GPT-4.1	Coding, instruction following, long docs	\$2.00	\$8.00	1M tokens
GPT-4.1 mini	Cost-efficient, high-volume tasks	\$0.40	\$1.60	1M tokens
GPT-4o mini	Fast, cheap, simple tasks	\$0.15	\$0.60	128K tokens
o3	Complex reasoning, maths, science	\$0.40	\$1.60	200K tokens
o4-mini	Fast reasoning, coding, visual tasks	\$1.10	\$4.40	200K tokens

When to use OpenAI: Start here. Best ecosystem, best docs, widest community support. Use GPT-4o for general tasks. Use o3 when you need reliable multi-step reasoning. Use GPT-4.1 mini or GPT-4o mini to keep costs low on high-volume calls.

Who they are

Anthropic was founded by ex-OpenAI researchers with a focus on AI safety. Their Claude models are widely regarded as the best at **following complex instructions precisely**, writing nuanced long-form content, and coding. Claude has a 200K token context window (vs 128K for standard GPT-4o), making it ideal for working with long documents. They pioneered **prompt caching** — cache your system prompt and pay up to 90% less on repeated calls.

Model lineup & pricing

Model	Best for	Input / 1M tokens	Output / 1M tokens	Context window
Claude Opus 4.7	Most capable, complex reasoning	\$5.00	\$25.00	200K tokens
Claude Opus 4.6	Highly capable, strong coding	\$5.00	\$25.00	200K tokens

Model	Best for	Input / 1M tokens	Output / 1M tokens	Context window
Claude Sonnet 4.6	<i>Best balance: quality vs cost</i>	\$3.00	\$15.00	200K tokens
Claude Haiku 4.5	<i>Fastest, cheapest, lightweight tasks</i>	\$1.00	\$5.00	200K tokens

When to use Anthropic: Use Claude Sonnet 4.6 as your primary model for agentic tasks — it follows complex multi-step instructions better than most. Use Haiku 4.5 for high-volume, low-complexity steps inside an agent. Enable prompt caching on long system prompts to cut costs dramatically. The 200K context window is a genuine advantage for long-document RAG.

Who they are

Google DeepMind builds the Gemini family. Their most unique advantage is scale: Gemini 2.5 Pro supports a **1 million token context window** — enough to fit entire codebases or thousands of pages. They are also the leaders in **multimodal AI**: a single API call can process text, images, video, and audio together. Gemini 2.5 Flash is one of the cheapest capable models on the market at \$0.05/1M input tokens, making it excellent for cost-sensitive production workloads.

Model lineup & pricing

Model	Best for	Input / 1M tokens	Output / 1M tokens	Context window
Gemini 2.5 Pro	Complex reasoning, massive context	\$1.25	\$10.00	1M tokens
Gemini 2.5 Flash	Speed + cost balance, multimodal	\$0.05	\$0.20	1M tokens
Gemini 3.1 Flash Lite	Ultra-cheap, high-volume tasks	\$0.05	\$0.20	1M tokens

When to use Google: Reach for Gemini when context length is your constraint — nobody else offers 1M tokens reliably. Gemini 2.5 Flash is the go-to model when you need a capable model at very low cost. Use Gemini when your input includes images, video, or audio alongside text. Access via Google AI Studio (free tier) or the Gemini API.

Who they are

Meta releases the **Llama** family of models as open-source weights under a permissive licence. This means you download the model weights and run them yourself — on your own GPU, a cloud VM, or via a third-party inference host. Meta does **not** operate a first-party API. The key advantages are: no per-token costs when self-hosted, full data privacy (nothing leaves your infrastructure), and the freedom to fine-tune the model on your own data.

Model lineup & pricing (via third-party hosts)

Model	Best for	Input / 1M tokens	Output / 1M tokens	Context window
Llama 3.3 70B	Best quality open-source model	~\$0.23	~\$0.40	128K tokens
Llama 3.1 8B	Lightweight, fast, cost-efficient	~\$0.02	~\$0.05	128K tokens
Llama 3.2 3B	On-device / edge deployment	~\$0.06	~\$0.06	128K tokens

Third-party hosts for Llama: **Together AI**, **Fireworks AI**, **Groq** (fastest), **AWS Bedrock**, **Azure AI**, **Deepinfra**. Prices vary between hosts — always compare.

When to use Meta / Llama: When data privacy matters (healthcare, finance) and you cannot send data to a third-party API. When you want to fine-tune on proprietary data. When you need to eliminate per-token API costs at scale. Llama 3.3 70B is competitive with GPT-4o mini in quality at a fraction of the cost via hosted providers.

Mistral Open + closed API

Flagship: Mistral Large 3 · Medium 3 · Devstral 2

Who they are

Mistral AI is a French startup that punches above its weight. They release smaller open-source models (like Mistral 7B) while also offering a closed API with their flagship large models. Their flagship **Mistral Large 3** is a sparse mixture-of-experts model with 675B total parameters (41B active) — delivering frontier-quality output at mid-tier pricing. They are particularly strong at **coding** (Devstral, Codestral) and are a strong choice if you are building in Europe and need GDPR-compliant AI infrastructure.

Model lineup & pricing

Model	Best for	Input / 1M tokens	Output / 1M tokens	Context window
Mistral Large 3	Frontier reasoning, complex tasks	~\$2.00	~\$6.00	128K tokens
Mistral Medium 3	Best quality-to-cost ratio in lineup	\$0.40	\$2.00	128K tokens
Mistral Small 4	Fast, general-purpose, cost-efficient	\$0.10	\$0.30	32K tokens
Devstral 2	Coding agent, repo-level tasks	\$0.40	\$2.00	128K tokens
Codestral	Code completion, fill-in-the-middle	\$0.30	\$0.90	256K tokens

When to use Mistral: Mistral Medium 3 offers excellent quality at \$0.40/1M input — one of the best cost-efficiency ratios in the market. Use Devstral 2 or Codestral for AI coding tasks. Choose Mistral if you need a GDPR-friendly European provider. Their open-source smaller models (Mistral 7B) are great for local experimentation.

Quick Comparison — All Five Providers

Provider	Flagship model	Price range (input)	Context window	Open-source?	Best known for
OpenAI	GPT-4o, GPT-4.1, o3	\$0.15 – \$2.50/1M	128K – 1M	No	Ecosystem, tooling, reasoning (o-series)
Anthropic	Claude Sonnet 4.6	\$1.00 – \$5.00/1M	200K	No	Instruction following, coding, safety
Google	Gemini 2.5 Pro/Flash	\$0.05 – \$1.25/1M	Up to 1M	No	Massive context, multimodal, low cost
Meta	Llama 3.3 70B	\$0.02 – \$0.40/1M	128K	Yes	Privacy, self-hosting, fine-tuning
Mistral	Mistral Large 3	\$0.10 – \$2.00/1M	32K – 256K	Partial	Cost-efficiency, coding, GDPR-friendly

How to Choose a Provider

If your situation is...
■ Starting a new project? → Default to OpenAI (GPT-4o or GPT-4.1). Best docs, widest support, easiest to debug.
■ Building an agent or complex pipeline? → Use Anthropic Claude Sonnet 4.6. Best instruction-following; 200K context window is an advantage.
■ Need the cheapest capable model? → Gemini 2.5 Flash at \$0.05/1M input. Or GPT-4o mini at \$0.15/1M.
■ Handling very long documents (>128K)? → Gemini 2.5 Pro or Anthropic Claude (200K context). Gemini goes up to 1M tokens.
■ Data privacy is a hard requirement? → Self-host Llama 3.3 70B. Nothing leaves your infrastructure.
■ Need to fine-tune on your own data? → Start with Llama 3 (open weights) or use OpenAI fine-tuning API.
■ Building a coding assistant? → Try Mistral's Devstral 2 or Codestral — purpose-built for code.
■ Operating in Europe / GDPR concerns? → Mistral AI is a European company with EU-based infrastructure.

Your Task Checklist

You have completed this task when all 5 items are done:
■ Create accounts — Sign up on OpenAI Platform (platform.openai.com) and Anthropic Console (console.anthropic.com). Get API keys for both. Store them in a .env file.

■ **Note the pricing** — Write down input and output prices for: GPT-4o, GPT-4.1 mini, Claude Sonnet 4.6, Claude Haiku 4.5, and Gemini 2.5 Flash. These five you will use most.

■ **Calculate a cost** — Estimate the cost of 10,000 API calls, each with a 1,000-token prompt and a 500-token response, using GPT-4o vs Claude Haiku 4.5. Which is cheaper and by how much?

■ **Explore Meta Llama** — Go to together.ai or fireworks.ai and find Llama 3.3 70B. Note its price per 1M tokens and compare to GPT-4o mini.

■ **Map provider to use case** — Write a one-line answer to: which provider would you use for (1) a customer support chatbot that handles 1M messages/day, (2) a legal document analyser needing 300K context, and (3) a coding assistant plugin?